

Maximum Likelihood Estimation

Mr Fugu Data Science



Youtube (<https://www.youtube.com/channel/UCbni-TDI-Ub8VIGaP8HLTNw/>) | Github (<https://github.com/MrFuguDataScience>)

Purpose & Outcome:

- ddfdf

Maximum Likelihood Estimation : allows us to estimate our "*unknown*" parameters for a given probability distribution. Think of it like this: we are trying to find the likelihood "*probability*" of generating this dataset.

- The parameters are found where they maximize the likelihood of generating our model.
- We are trying to figure out what model best describes our data. Having domain knowledge can help, in these circumstances.

Desireable Effects : (if you have a large sample size)

- As the sample size increases, we reach an unbiased variance estimator
- The ability to generate a normal distribution and confidence intervals can be used with large enough samples
- The ability to form hypothesis testing

Drawbacks :

- If you have small sample size there will be bias
- Can be a computational hog

Each probability distribution will have unique parameters. As the parameters change, a different probability distribution is generated.

- The parameters for each model will fall into a family of distributions based on the parameters you have. The parameters here will be (μ, σ)

For example you may have a normal distribution, but the mean (μ) and standard deviation (σ) are not the same from one model to the other.

- Therefore, you may have a family of normal distributions with varying parameters.

Plot showing different normals and explain as family of functions

- These two plots reflect the parameters of a Normal Distribution you may want to estimate. If this were a different problem you would evaluate other parameters and which may not be $[\mu, \sigma]$

```

In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn') # pretty matplotlib plots
plt.rcParams['figure.figsize'] = (12, 8)

x = np.linspace(-5, 5, 5000)
mu = [0,0,0]
sigma = [.5,1,3]

y_=[]
for i in range(len(sigma)):

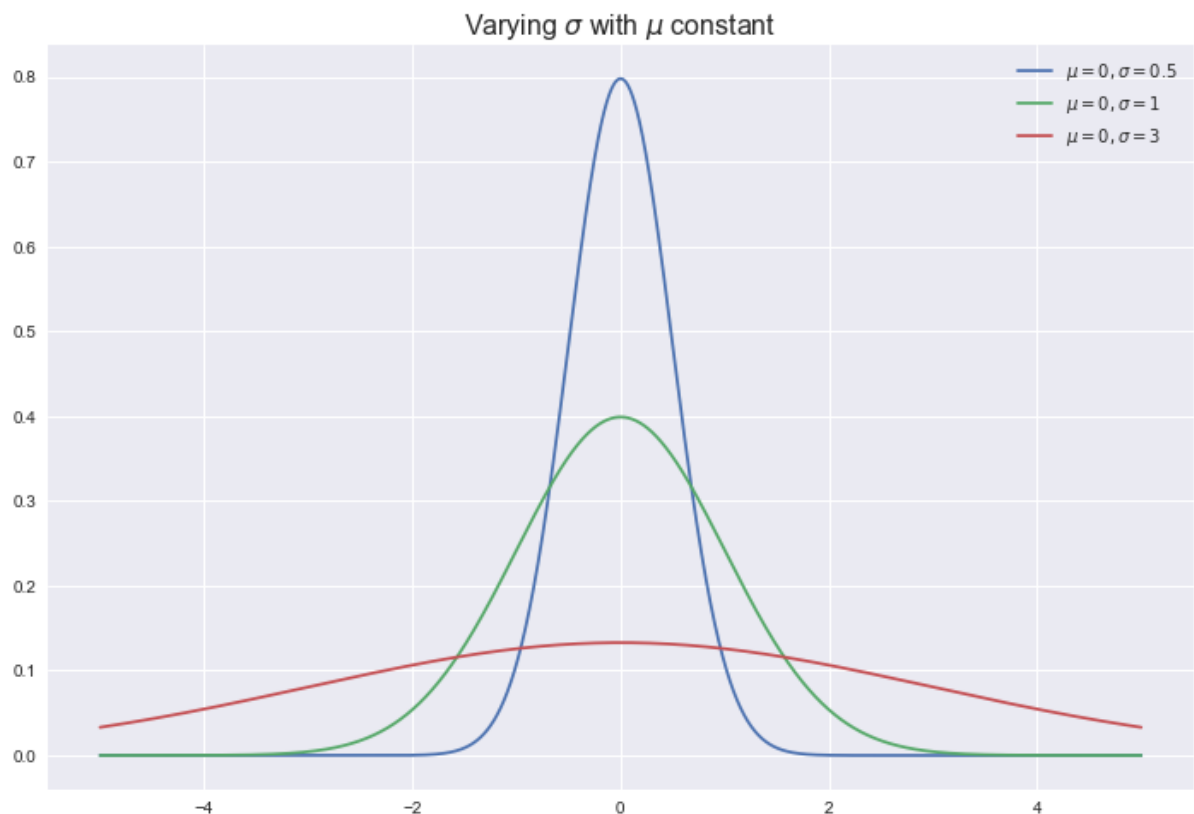
    y = (1 / (np.sqrt(2 * np.pi * np.power(sigma[i], 2)))) * \
        (np.power(np.e, -(np.power((x - mu[i]), 2) / (2 * np.power(sigma[i],
2))))))

    y_.append(y)

labels=[('$\mu=0,\sigma=0.5$'),('$\mu=0,\sigma=1$'),('$\mu=0,\sigma=3$'
)]

for i in range(len(y_)):
    plt.plot(x,y_[i],label=labels[i])
    plt.title('Varying $\sigma$ with $\mu$ constant', size=16)
    plt.legend();

```



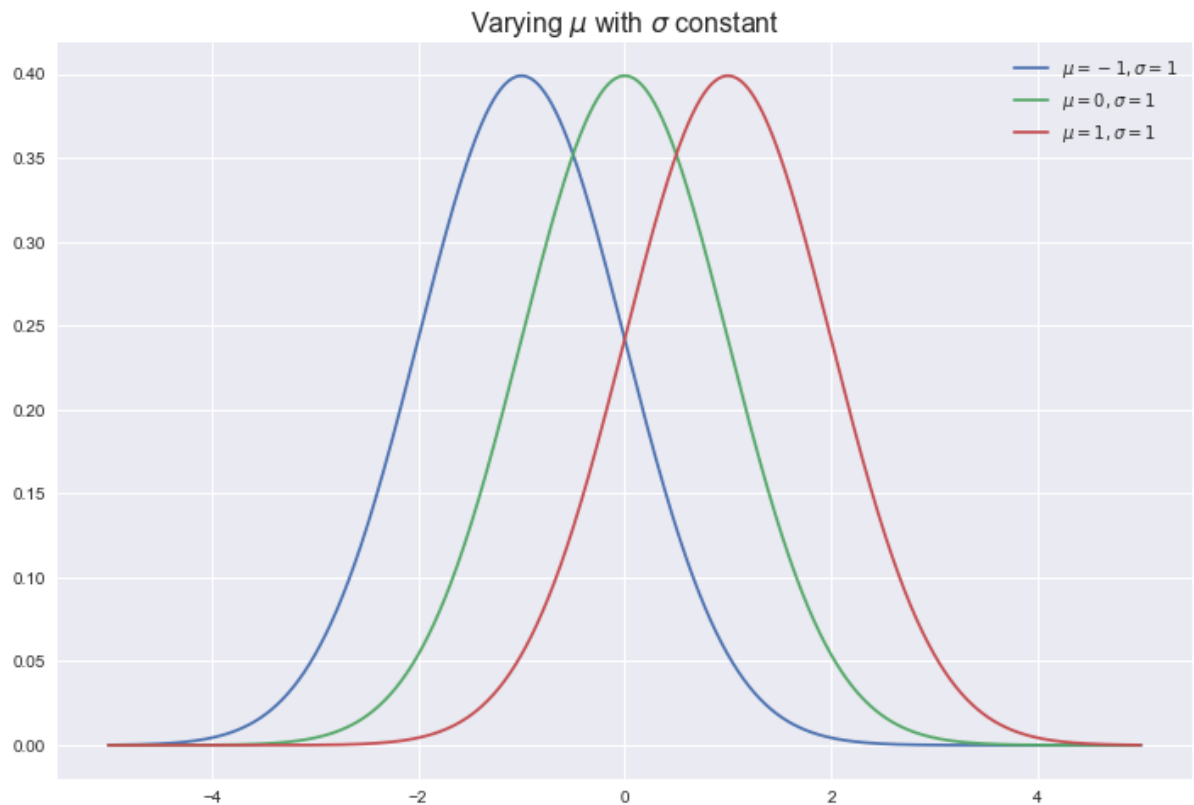
```

In [2]: x = np.linspace(-5, 5, 5000)
mu_1=[-1,0,1]
sigma_1=[1,1,1]
y_1=[]
for i in range(len(sigma)):

    y_01= (1 / (np.sqrt(2 * np.pi * np.power(sigma_1[i], 2)))) * \
        (np.power(np.e, -(np.power((x - mu_1[i]), 2) / (2 * np.power(sigma_1
[i], 2)))))
    y_1.append(y_01)

labels_1=[('$\mu=-1,\sigma=1$'),('$\mu=0,\sigma=1$'),('$\mu=1,\sigma=1$'
)]
for i in range(len(y_)):
    plt.plot(x,y_1[i],label=labels_1[i])
    plt.legend();
    plt.title('Varying $\mu$ with $\sigma$ constant',size=16)

```



In the real world we are given some data that we observed, and need to find some model of interest that will best describe our data using a PDF that will most likely represent our data.

- Realistically, we do not always have the option of transforming our data (*log, reciprocal, square root, etc*) to save the day.
-

Instead we have: $L(\theta; x_i) = f(x; \theta)$

$L(\theta; x_i)$: This means that we have the likelihood of some true parameter θ given the observed data x

well what is the right side then? : it is our PDF and there is a distinction

$f(x; \theta)$: is a function of data given a particular set of parameters

I am using the semicolon to denote parameters with respect to x instead of confusing conditional probability at this point, which would have the vertical bars

Math of MLE:

Nasty Truth:

- **First**, there doesn't have to be a unique value of the MLE
 - **Second**, there may not be a value that exists for the MLE

Hint: because you need to think about concavity, this is a problem if we are dealing with machine learning algorithms and a bit topic.

Super Fun Happy Times, wait what?

ಠ_ಠ?

MaTh TiMe Y'aLL :

$$L(\theta; x_i) = f(x; \theta)$$

$$f(x; \theta) = \prod f(X|\theta) =$$

$$f(x_1, x_2, x_3 \dots x_n | \theta) =$$

$$f(x_1 | \theta) * f(x_1 | \theta) * f(x_2 | \theta) * \dots f(x_n | \theta)$$

But, now the trickery $\backslash_(_)_/$:

We will decide to do the Log because, it will get rid of our multiplications and simply everything. Also, since the log is strictly increasing it shouldn't effect our function.

- We will need to find the Max/Min, hence we will need the first derivative with respect to θ and equate to zero.

$$\log(L(\theta; x_i)) = \log[f(\theta|x_1) * f(\theta|x_1) * f(\theta|x_2) * \dots f(\theta|x_n)]$$

$$= \log(f(\theta|x_1)) + \log(f(\theta|x_1)) + \log(f(\theta|x_2)) + \dots \log(f(\theta|x_n))$$

$$= \sum_{i=1}^n \log(f(\theta|x))$$

Derivative Time

$$\text{First Derivative: } \frac{\partial \ln(L(\theta|x))}{\partial \theta_i} = 0$$

$$\text{Second Derivative: } \frac{\partial^2 \ln(L(\theta|x))}{\partial \theta_i^2} < 0$$

At this point if you have a suspected distribution, you would replace with your distribution function. There are times when:

Real World : Either the data are highly paramertized , non-convex (*are a few reasons*) and you will need to use optimazation techinues to attempt solving these problems with approximation methods.

- Examples : *Newton's method, Fisher scoring, various conjugate gradient-based approaches, steepest descent, Nelder-Mead type (simplex) approaches, BFGS* and a wide variety of other techniques (from: <https://www.analyticsvidhya.com/blog/2018/07/introductory-guide-maximum-likelihood-estimation-case-study-r/> (<https://www.analyticsvidhya.com/blog/2018/07/introductory-guide-maximum-likelihood-estimation-case-study-r/>))

a great resource: <https://www.stat.tamu.edu/~suhasini/teaching613/chapter2.pdf>
(<https://www.stat.tamu.edu/~suhasini/teaching613/chapter2.pdf>)

Assume we have some (* fake*) data based on the preceived IQ for people who vote for a particular president

```
In [122]: # Histogram with density plot

import seaborn as sns
import matplotlib.mlab as mlab
%matplotlib inline

fake_data=[30,40,62,61,85,80,16,55,26,53,51,70,116,
           66,51,101,99,88,71,55,51,35,62,71,81,44]



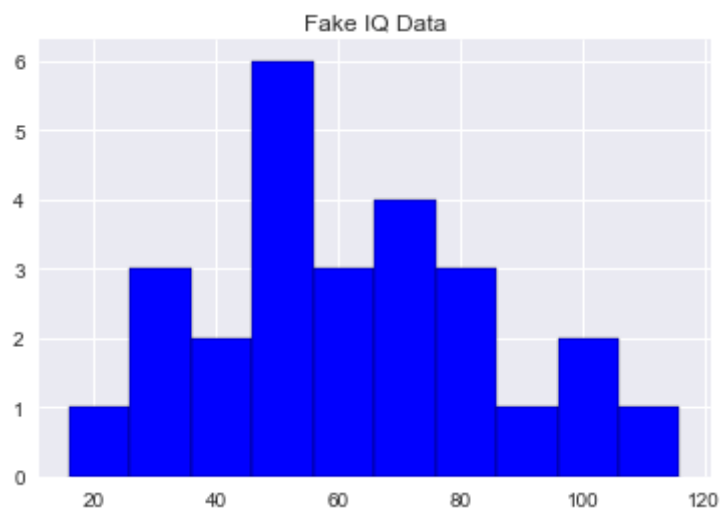
plt.hist(fake_data, color = 'blue', edgecolor = 'black')
plt.title('Fake IQ Data')
plt.subplots()

plt.hist(np.log(fake_data), color = 'green', edgecolor = 'black')
plt.title('Fake IQ Data (Log Transformed)')
plt.subplots()

plt.hist(np.sqrt(fake_data), color = 'red', edgecolor = 'black')
plt.title('Fake IQ Data (Sqrt Transformed)')
# plt.subplot()
# plt.hist(np.log(fake_data), color = 'green', edgecolor = 'black')
# plt.title('Fake IQ Data (Log Transformed)')
```



```
Out[122]: Text(0.5, 1.0, 'Fake IQ Data (Sqrt Transformed)')
```



Lets look at this a little more clearly:

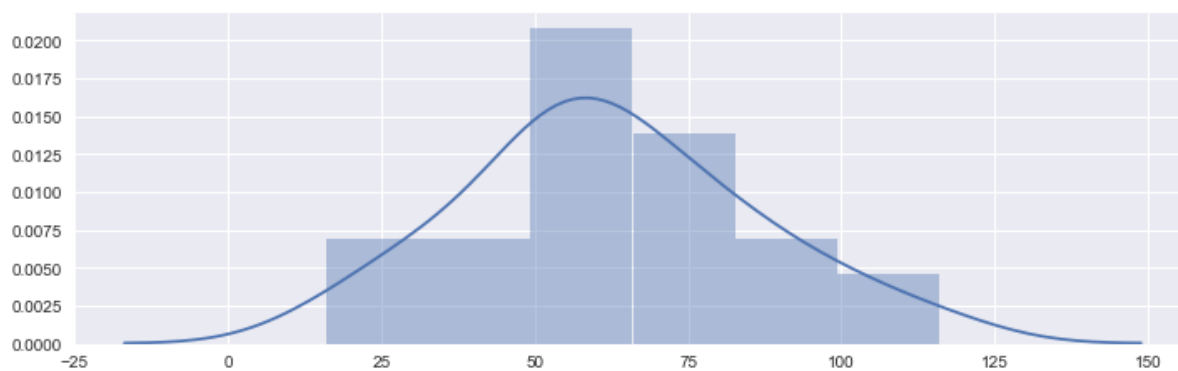
```
In [296]: fig, axarr = plt.subplots(2, 1, figsize=(12, 8))

sns.distplot(fake_data, ax=axarr[0])
plt.suptitle('Fake IQ Data: [original]', size=14)


sns.distplot(np.sqrt(fake_data), ax=axarr[1])
plt.title('Fake IQ Data: [sqrt transformed]', size=14)

plt.show()
```

Fake IQ Data: [original]



Fake IQ Data: [sqrt transformed]



transform data and looks close? What now?

- Well, lets consider that maybe we think these data reflect a normal distribution.
 - We could do a comparison after finding the parameters and see how it matches up.
 - We also, may end up in a situation where we need to reconsider our distribution as well. (*which is*) very likely in the real world.

For simplicity we are assuming Normal distribution :

$$f(x_1, x_2, x_3 \dots x_n; \mu, \sigma) =$$

Replace with distribution:

$$P(X; \mu, \sigma) = \frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) =$$

$$\prod \frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) =$$

Expand:

$$\frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_1-\mu)^2}{2\sigma^2}\right) * \frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_2-\mu)^2}{2\sigma^2}\right) * \frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_3-\mu)^2}{2\sigma^2}\right) * \dots * \frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_n-\mu)^2}{2\sigma^2}\right)$$

Now take the Log:

$$\ln(L(\theta; X)) = \ln\left(\frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_1-\mu)^2}{2\sigma^2}\right)\right) + \ln\left(\frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_2-\mu)^2}{2\sigma^2}\right)\right) + \ln\left(\frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_3-\mu)^2}{2\sigma^2}\right)\right) + \dots + \ln\left(\frac{1}{\sigma\sqrt{(2\pi)}} \exp\left(-\frac{(x_n-\mu)^2}{2\sigma^2}\right)\right)$$

Simplify :

$$\ln(L(\theta; X)) = \ln\left(\frac{1}{\sigma\sqrt{(2\pi)}}\right) + \ln\left(\exp\left(-\frac{(x_1-\mu)^2}{2\sigma^2}\right)\right), \text{ this is done for each } x_i \text{ so I will simplify:}$$

$$= \frac{-1}{2} \ln(2\sigma\pi) - \frac{(x_1-\mu)^2}{2\sigma^2}$$

$$= \frac{-1}{2} \ln(2\pi) - \frac{1}{2} \ln(\sigma) - \frac{(x_1-\mu)^2}{2\sigma^2}, \text{ Now repeat for a } X=[x_1, x_2 \dots x_n]$$

To speed this up and understand that we have: (n) observations (data points)

$$= \frac{-n}{2} \ln(2\pi) - n * \ln(\sigma) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \text{ This is what we will take the derivative of}$$

Now the Derivatives :

The first derivative will be the slope of the log-likelihood curve:

$$\frac{\partial L}{\partial \mu} = 0 - 0 - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 =$$

$$\frac{1}{\sigma^2} \sum_{i=1}^n x_i - \mu =$$

$$\frac{1}{\sigma^2} \sum_{i=1}^n x_i - n\mu$$

$$\hat{\mu} = \frac{\sum_{i=1}^n x_i}{n} = \text{sample mean}$$

-----Now Standard Deviation-----

short cut:

$$= -\frac{n}{2} \ln(2\pi) - n * \ln(\sigma) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

$$\frac{\partial L}{\partial \sigma} = 0 - \frac{n}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^n (x_i - \mu)^2 =, \text{ set equal to zero}$$

$$-\frac{n}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^n (x_i - \mu)^2 = 0$$

Manipulate:

$$\frac{1}{\sigma^3} \sum_{i=1}^n (x_i - \mu)^2 = \frac{n}{\sigma}$$

$$\hat{\sigma} = \sqrt{\frac{\sum_{i=1}^n (x_i - \mu)^2}{n} = \frac{n}{\sigma}}, \text{ biased estimator}$$

side note: the (n-1) will be unbiased look into proof, second webpage below

<https://machinelearningmastery.com/probability-density-estimation/>


(<https://machinelearningmastery.com/probability-density-estimation/>)

<https://daijiang.name/en/2014/10/08/mle-normal-distribution/> (<https://daijiang.name/en/2014/10/08/mle-normal-distribution/>)

Check the CDF: so we can get an idea of if these data will fit a Normal Distr

```
In [511]: kwargs = {'cumulative': True}
# sns.distplot(fake_data, hist_kws=kwargs, kde_kws=kwargs)
sns.kdeplot(fake_data, cumulative=True)
```

```
Out[511]: <matplotlib.axes._subplots.AxesSubplot at 0x1a424b5208>
```



We can directly compute the μ & σ from the data, but this doesn't mean that we are actually representing the population the data came from.

- Hence, the use of MLE .

```

In [452]: print(np.mean(fake_data),np.std(fake_data))
fig_dims = (10, 6)
fig, ax = plt.subplots(figsize=fig_dims)
sns.distplot(np.random.normal(np.mean(fake_data),np.std(fake_data),1000
))
# plt.legend('MLE')
sns.distplot(fake_data,color='red')
plt.axvline(np.mean(fake_data),color='b', linestyle='--')
plt.legend(['Mean original data','Fitted MLE','Orig. Data'],bbox_to_ancho
r=(1, 1),
           loc=2, borderaxespad=0.) #,fontsize=25
plt.show()

# sigma=np.std(fake_data)
# mu=np.mean(fake_data)

# vals=[]
# for x in sorted(fake_data):
#     norm_d=1/(sigma*np.sqrt(2*np.pi))*np.exp(-(x-mu)^2/2*sigma^2)
#     norm_d=1/(sigma*np.sqrt(2*np.pi))*np.exp(-np.power(x - mu, 2.) /
#     (2 * np.power(sigma, 2.)))
#     vals.append(norm_d)
# sorted(fake_data)

```

62.30769230769231 23.57601328398981



```


In [508]: import numpy as np
from scipy.stats import norm
import matplotlib.pyplot as plt
import scipy.stats as stats
fig_dims = (14, 8)
fig, ax = plt.subplots(figsize=fig_dims)
mean_ = np.mean(fake_data)
sd_ = np.std(fake_data)

pdf = stats.norm.pdf(sorted(fake_data), mean_, sd_)
plt.plot(sorted(fake_data), pdf, color='red') # including h here is crucial

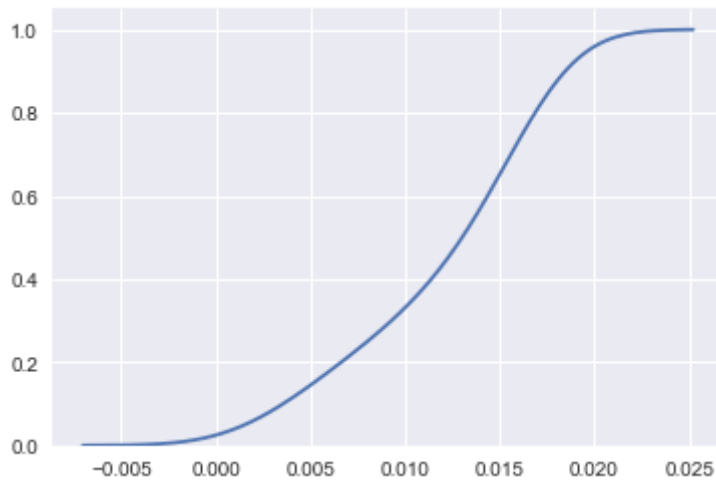
sns.distplot(fake_data, color='green')
plt.axvline(np.mean(fake_data), color='b', linestyle='--')

plt.legend(['Mean original data', 'Fitted MLE', 'Orig. Data'], bbox_to_anchor=(1, 1),
           loc=2, borderaxespad=0., fontsize=11)
plt.show()

```



```
In [513]: sns.kdeplot(pdf,cumulative=True)  
plt.show()
```



Last Thoughts:

- You can't always, assume your data will fit an exact model
- Approximation methods are usually a good call on real data
- Do not blindly, believe everything you hear!

SHARE, Share &

SUBscribe

Citations & Help:



<https://www.itl.nist.gov/div898/handbook/apr/section4/apr412.htm>
(<https://www.itl.nist.gov/div898/handbook/apr/section4/apr412.htm>)

https://en.wikipedia.org/wiki/Maximum_likelihood_estimation
(https://en.wikipedia.org/wiki/Maximum_likelihood_estimation)

<https://towardsdatascience.com/probability-concepts-explained-maximum-likelihood-estimation-c7b4342fdbb1>
(<https://towardsdatascience.com/probability-concepts-explained-maximum-likelihood-estimation-c7b4342fdbb1>)

<https://machinelearningmastery.com/what-is-maximum-likelihood-estimation-in-machine-learning/>
(<https://machinelearningmastery.com/what-is-maximum-likelihood-estimation-in-machine-learning/>)

<https://www.sciencedirect.com/topics/engineering/maximum-likelihood-estimation>
(<https://www.sciencedirect.com/topics/engineering/maximum-likelihood-estimation>)

<https://www.math.arizona.edu/~jwatkins/n-mle.pdf> (<https://www.math.arizona.edu/~jwatkins/n-mle.pdf>) (math background)

<https://medium.com/@rrfd/what-is-maximum-likelihood-estimation-examples-in-python-791153818030>
(<https://medium.com/@rrfd/what-is-maximum-likelihood-estimation-examples-in-python-791153818030>)

<https://python.quantecon.org/mle.html> (<https://python.quantecon.org/mle.html>) (python example)

<http://times.cs.uiuc.edu/course/410/note/mle.pdf> (<http://times.cs.uiuc.edu/course/410/note/mle.pdf>) (good paper)

https://emredjan.github.io/blog/2017/07/19/plotting-distributions/?source=post_page-----c5ebaafdeedd-----
(https://emredjan.github.io/blog/2017/07/19/plotting-distributions/?source=post_page-----c5ebaafdeedd-----)

https://www2.math.ethz.ch/education/bachelor/lectures/fs2016/math/wr_s/solution_11.pdf
(https://www2.math.ethz.ch/education/bachelor/lectures/fs2016/math/wr_s/solution_11.pdf) (great read)

https://psu-psychology.github.io/psy-597-SEM/04_latent_variables_estimation/maximum_likelihood.html
(https://psu-psychology.github.io/psy-597-SEM/04_latent_variables_estimation/maximum_likelihood.html)

http://web.vu.lt/mif/a.buteikis/wp-content/uploads/PE_Book/3-4-UnivarMLE.html
(http://web.vu.lt/mif/a.buteikis/wp-content/uploads/PE_Book/3-4-UnivarMLE.html)

<https://stackoverflow.com/questions/53128352/legend-overlapping-plot-area-in-seaborn>
(<https://stackoverflow.com/questions/53128352/legend-overlapping-plot-area-in-seaborn>)

<https://people.missouristate.edu/songfengzheng/Teaching/MTH541/Lecture%20notes/MLE.pdf>
(<https://people.missouristate.edu/songfengzheng/Teaching/MTH541/Lecture%20notes/MLE.pdf>)

<https://www.youtube.com/watch?v=ttmKa1DovfI> (<https://www.youtube.com/watch?v=ttmKa1DovfI>)

<https://moonbooks.org/Articles/How-to-calculate-a-log-likelihood-in-python-example-with-a-normal-distribution-/> (<https://moonbooks.org/Articles/How-to-calculate-a-log-likelihood-in-python-example-with-a-normal-distribution-/>)

<https://stackoverflow.com/questions/20011494/plot-normal-distribution-with-matplotlib/20026448>
(<https://stackoverflow.com/questions/20011494/plot-normal-distribution-with-matplotlib/20026448>)

<https://stackoverflow.com/questions/53128352/legend-overlapping-plot-area-in-seaborn>
(<https://stackoverflow.com/questions/53128352/legend-overlapping-plot-area-in-seaborn>)

<https://stackoverflow.com/questions/51417483/mean-median-mode-lines-showing-only-in-last-graph-in-seaborn> (<https://stackoverflow.com/questions/51417483/mean-median-mode-lines-showing-only-in-last-graph-in-seaborn>)

http://web.vu.lt/mif/a.buteikis/wp-content/uploads/PE_Book/3-4-UnivarMLE.html
(http://web.vu.lt/mif/a.buteikis/wp-content/uploads/PE_Book/3-4-UnivarMLE.html)